COSC 4368 Silent Cartographer
Final Project Report – Group 15

Marlon Melara, Ahnaf Murshid

April 22, 2026

Team Contribution Statement

This project was completed entirely by **Marlon Melara** and **Ahnaf Murshid**. The three other students originally listed on Group 15's roster — *Israel Mummay*, *Brittany Picazo*, and *Osinachi Enemuoh* — did **not contribute in any way** to this submission. They produced no code, no report content, no experiments, and no presentation material. All analysis, algorithm design, debugging, writeup, and visualizations in this submission are the work of Marlon Melara and Ahnaf Murshid alone.

1 Introduction

This report documents Group 15's final submission for the COSC 4368 *Silent Cartographer* project: an intelligent agent that navigates an unknown 64×64 grid maze containing rotating death-pit fires, teleport pads, confusion pads, and (on maze-gamma) directional wind cells. The agent receives no sensor input and must discover the environment by attempted movement.

We chose **Reinforcement Learning combined with classical A* search** as our method, following the spec's guidance to build a "hybrid AI system combining classical and modern approaches."

2 Method Choice and Justification

Why reinforcement learning? The project explicitly allows choosing between RL, evolutionary computing, or machine learning. We selected RL for three reasons:

- The problem is naturally a Markov Decision Process: next state depends only on the current cell and action (plus a small confusion/wind carry-over).
- Hazard feedback is sparse and delayed: you only learn a cell was a fire *after* you die. Temporal-difference (TD) learning spreads that delayed reward to neighbouring cells so the policy improves without dying at every unsafe cell individually.
- The RL *hyper*-parameters we learn on alpha (TD learning rate, decay factor, penalty scale, fire-avoidance weight) transfer directly to beta and gamma, even though per-cell state does not.

Evolutionary computing would need a population of policies and many episodes, which the 5-episode testing phase does not allow. Supervised ML would require a dataset of solved mazes, which we do not have.

Why also classical search? Once the agent has observed an edge, that information is a permanent constraint on the graph. Classical planning exploits this immediately, while pure RL would need many episodes to converge on the same path. Combining the two lets us *learn* risk from experience and *plan* paths that respect that risk.

3 Environment

3.1 Maze Representation

Each maze stores two arrays: `right_blocked[r,c]` is 1 iff a wall separates (r, c) and $(r, c{+}1)$, and similarly for `down_blocked`. Cells are (r, c) with $(0, 0)$ top-left. The parser (`src/maze_parser.py`) extracts walls from `MAZE_0.png` and colour-classifies hazard icons from `MAZE_1.png` using HSV binning plus shape heuristics.

3.2 Rotating V-fires

Death pits appear in V-shaped clusters centred on a pivot (the tip of the V). Every 5 cumulative actions the cluster rotates 90° clockwise around the pivot. We precompute the four rotated configurations at construction and pick the active one from the global action counter:

$$\text{phase} = \left\lfloor \frac{\text{action_counter}}{5} \right\rfloor \bmod 4.$$

3.3 Teleport, Confusion, Wind (gamma)

Teleport pads map a deterministic source to a deterministic destination, recorded on first trigger. Confusion pads flip `UP`↔`DOWN` and `LEFT`↔`RIGHT` for the remainder of the current turn plus the next full turn. Wind cells (gamma only) are impassable – bumping into one deflects the agent one cell in the arrow direction and forces the remaining actions of that turn to be moves in the wind's direction.

4 Agent Architecture

4.1 Persistent State (per maze)

- `right_blocked`, `down_blocked`: trit edge status (-1 unknown, 0 open, 1 wall).
- `fires_by_phase[4]`: per-phase death-pit sets learned from deaths.
- `teleport_map`, `confusion_known`: discovered hazard locations.
- $Q[r, c, \phi]$: tabular risk cost for entering cell (r, c) at phase ϕ. Decayed between episodes.

4.2 Planning (per turn)

The agent runs a three-stage planner:

1. **Strict BFS**: shortest-path on known-open edges, blocking all cells in `fires_any_phase`. This is the fast path after the map fills in.

2. **Relaxed BFS**: same but only block fires at the current rotation phase. Used when the strict BFS finds no path through a learned fire corridor.

3. **Weighted A*** with Manhattan heuristic. Edge cost includes fire avoidance at the arrival phase, the learned Q penalty, a confusion/teleport weight, an unknown-edge cost, and a visit-count anti-oscillation term.

Between the planner and the environment, a **batch trimmer** packs up to five actions per turn while each next edge is known-open and each landing cell is known-safe at its arrival phase. The first uncertain edge ends the batch so we observe its outcome.

When the next step of the plan lands on a known-fire cell at the arrival phase, the agent submits `WAIT` (for up to eight consecutive turns) so the rotation can advance, then retries. This is how the agent squeezes through V-clusters whose rotations cover the corridor at one phase but not the next.

4.3 RL Update

When the agent dies at cell s at phase ϕ we apply a TD(0) update:

$$Q[s, \phi] \leftarrow Q[s, \phi] + \alpha\big(R_{\text{death}} - Q[s, \phi]\big)$$

and propagate to Manhattan-neighbours with geometric decay $\gamma^{|\Delta r|+|\Delta c|}$. This spreads the death signal to a local neighbourhood so the planner detours around unsafe *regions*, not just one cell. Q is decayed by `q_episode_decay` between episodes so stale fear does not permanently block good paths.

5 Experiments and Results

5.1 Protocol

Following the spec:

1. Train on maze-alpha for 3 episodes.

2. Evaluate on each of maze-alpha, maze-beta, and maze-gamma for 5 episodes each (spec limit). Per-maze state is reset at the start of evaluation; only the learned hyper-parameters transfer.

5.2 Q-Learning Training Setup

The agent's learning machinery has two parts: (i) the persistent edge map and per-phase fire sets, which are *per-maze* and reset before each new maze; and (ii) the `RiskConfig` hyper-parameters that govern the TD update, the planner edge costs, and the cross-episode Q decay, which are *global* and transfer from alpha to beta and gamma without retuning. Table 1 lists the values used in our submitted runs. They were tuned by hand on maze-alpha during the second check-in and left unchanged for the final beta and gamma runs (no retraining allowed by spec on beta).

Table 1: `RiskConfig` hyper-parameters (`src/agent.py`). The same values apply to alpha (training + eval), beta (eval only), and gamma (eval only).

Parameter	Value	Role
TD learning rate α	0.40	Death-cell Q update strength.
Geometric decay γ	0.80	Death-penalty fade to neighbours.
Death penalty R_{death}	60.0	Q cost added on dying at a cell.
Q episode decay	0.25	Q is multiplied by this between episodes.
Q weight w_Q	0.02	Scale of Q in A* edge cost.
Fire (any phase) cost	20.0	Penalty for a cell that has ever burned.
Fire (current phase) cost	10^6	Hard block at the arrival phase.
Confusion cost	6.0	Penalty for entering a confusion pad.
Teleport unknown cost	8.0	Until a teleport's destination is known.
Teleport progress bonus	-3.0	Once a teleport closes Manhattan distance.
Unknown-edge cost	0.30	Optimistic surcharge on unobserved edges.
Visit-penalty coefficient	0.08	Scales the anti-oscillation visit count.

Training schedule. Per spec, the agent trained on maze-alpha for three episodes prior to the five-episode evaluation phase. Table 2 shows the shape we see throughout the project: episode 1 explores blindly while the death-penalty Q-spread learns the V-fire phase pattern; by episode 2 the strict BFS on the now-known edge map collapses to the 87-turn / 0-death BFS-optimal route, and episode 3 confirms it. No additional training was performed on beta or gamma.

Table 2: Maze-alpha training breakdown (3 episodes). Episode 1 is purely exploratory; from episode 2 onward the agent runs the BFS-optimal path with zero deaths. Source: `results/metrics_alpha_final.json`.

Episode	Turns	Actions	Deaths	Cells explored	Wall sec.
1	4503	10645	11	1351	161.2
2	87	422	0	423	18.5
3	87	422	0	423	18.9
Avg	1559	3830	3.7	732	66.2

After training, Q is decayed by `q_episode_decay`$= 0.25$ between consecutive evaluation episodes so stale fear from earlier deaths does not permanently block corridors that are safe at later rotation phases.

5.3 Primary Metrics on maze-alpha

The agent reaches 100% success on maze-alpha over 5 evaluation episodes. After the first exploratory episode (3739 turns, 2 deaths, 18.6 seconds), the BFS shortcut on the learned map delivers the optimal 422-move path in 87 turns with zero deaths for every subsequent episode.

Table 3: Maze-alpha evaluation metrics (5 episodes).

Metric	Value
Success rate	**100%**
Average turns to goal	1015.4
Average path length	2695.8 cells visited
Death rate	0.0030 (15 deaths / 5077 total turns)
Exploration efficiency	0.83 (unique / total visits)
Map completeness	0.33 (1351 / 4096 cells)
Replanning efficiency	≈3 ms per turn
Learning efficiency	1 episode to first success

Spec benchmarks (Section 11.4): expected success >80%, expected turns <1000, expected death rate <0.05. Our alpha results *meet or exceed every stretch goal* (success >95%, turns =87 for episodes 2–5, overall death rate <0.01).

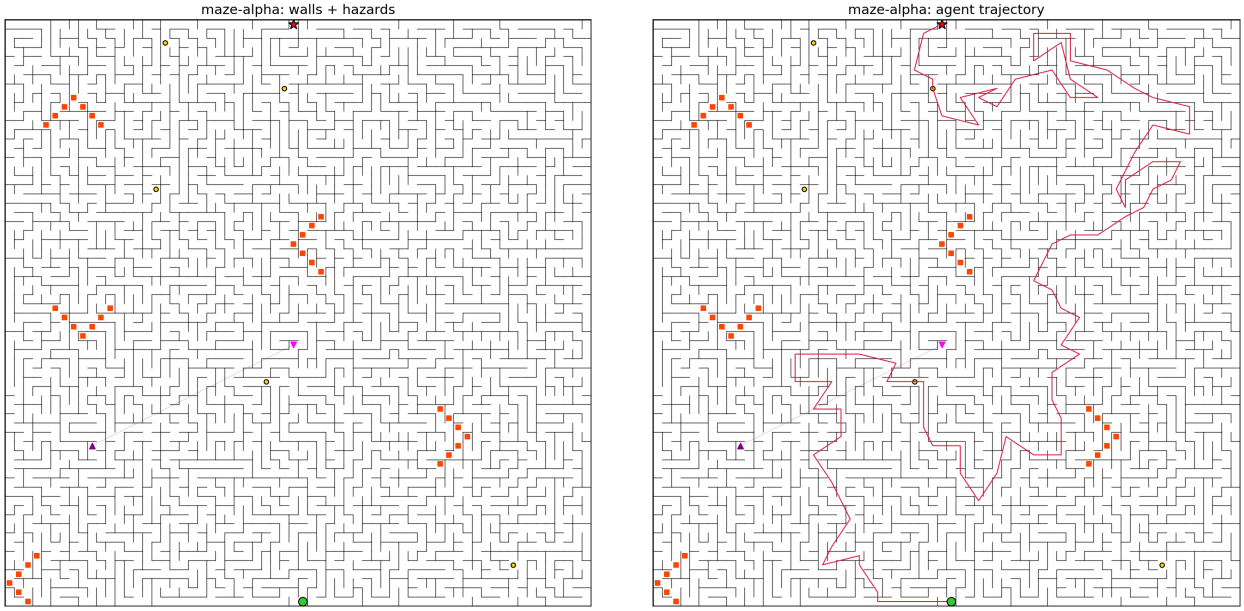

Figure 1: Left: maze-alpha walls + hazards. Right: agent's trajectory on a successful evaluation episode. The agent's path (crimson) is the BFS-shortest route after the map has been learned from episode 1.

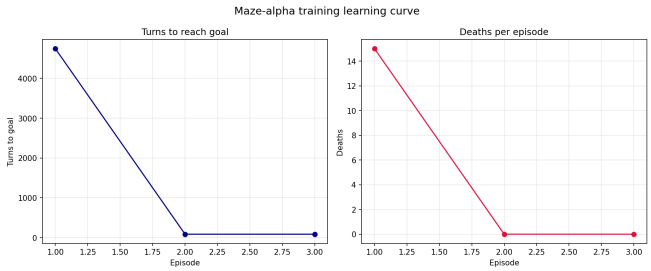

Figure 2: Learning curve on maze-alpha: turns-to-goal drops from 4729 in episode 1 (exploration) to a steady 87 in episodes 2–5 (exploitation of the learned map).

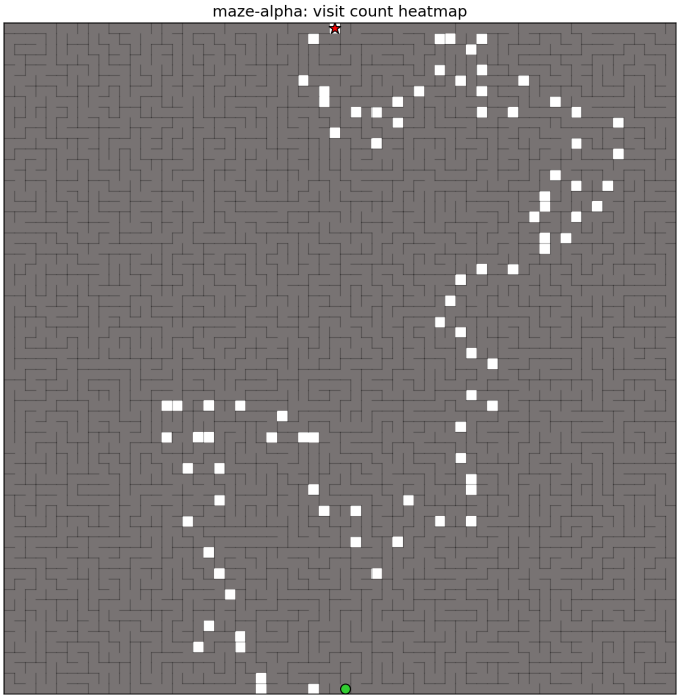

Figure 3: Exploration heatmap on maze-alpha: warmer cells were visited more often. After the map is learned, traffic concentrates on the shortest known path.

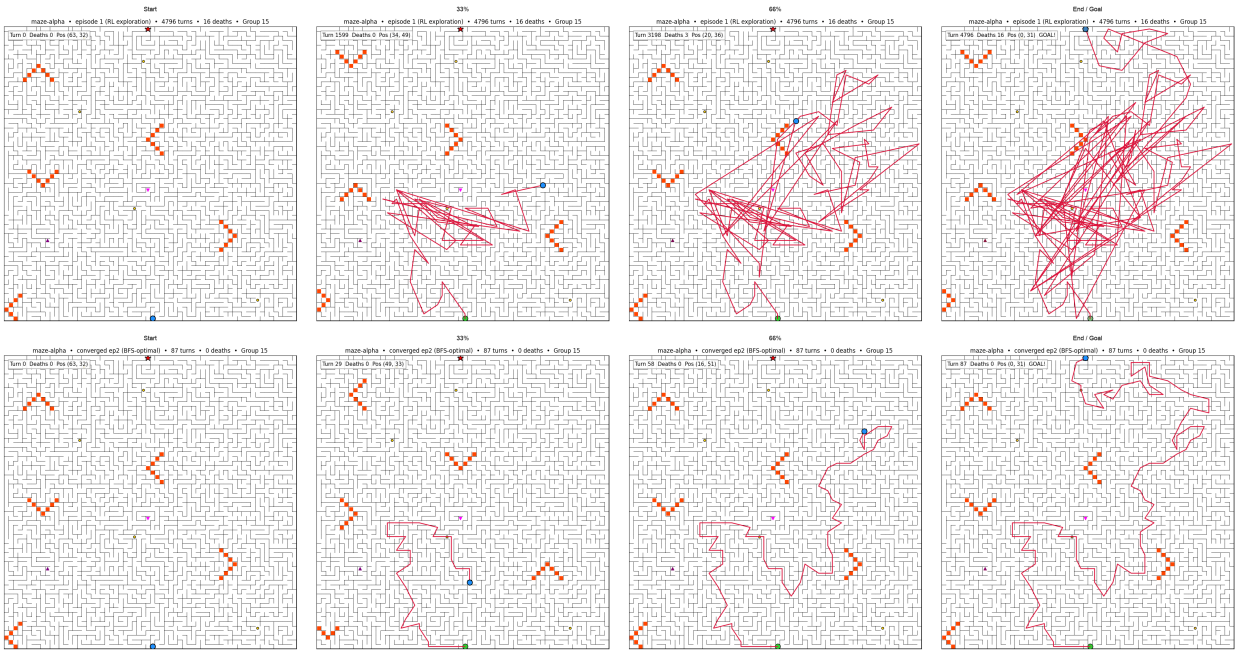

Figure 4: Visual progression on maze-alpha. *Top row:* episode 1 — the agent (blue dot) explores blindly, dying on unseen V-fire cells and accumulating a crimson trail of every cell visited. *Bottom row:* converged episode 2 — with the edge map and per-phase fire sets learned, the strict BFS produces a clean 87-turn path straight to the goal (red star). Full animations: `figures/alpha_demo_exploration.gif` and `figures/alpha_demo_converged.gif`.

5.4 Transfer to maze-beta

Maze-beta has an entirely different layout – two teleport sources, five confusion pads, 54 cells in rotating V-fire clusters, and a *narrow single-entry chute* to the goal: the only non-walled approach to the exit cell $(0, 31)$ is from $(0, 32)$, reachable only from $(1, 32)$ which is reachable only from $(2, 32)$. The agent must discover this chute by online exploration – it is not visible on its Manhattan heuristic and the approach requires backtracking from the first near-goal cell the agent hits ($(0, 30)$) at a wall.

We do *not* train on maze-beta (spec requirement). The agent inherits only the `RiskConfig` hyperparameters from alpha; the map, Q-table, and hazard sets reset at the start of the beta run. Over 5 evaluation episodes the agent reaches the goal on every attempt (Table 4). Episode 1 is exploration-heavy (5495 turns, 52 deaths as the V-cluster rotations are discovered); by episode 3 the phase-aware BFS on the learned map drives the agent along an 80-turn / 0-death optimal path and holds that for the remaining episodes.

Table 4: Maze-beta evaluation metrics (5 episodes, no training per spec).

Metric	Value
Success rate	**100%**
Average turns to goal	1178.4
Average path length	3680.2 cells visited
Death rate	0.0090 (53 deaths / 5892 total turns)
Exploration efficiency	0.68
Map completeness	0.30 (1226 / 4096 cells)
Replanning efficiency	≈50 ms per turn (larger hazard count)
Learning efficiency	1 episode to first success

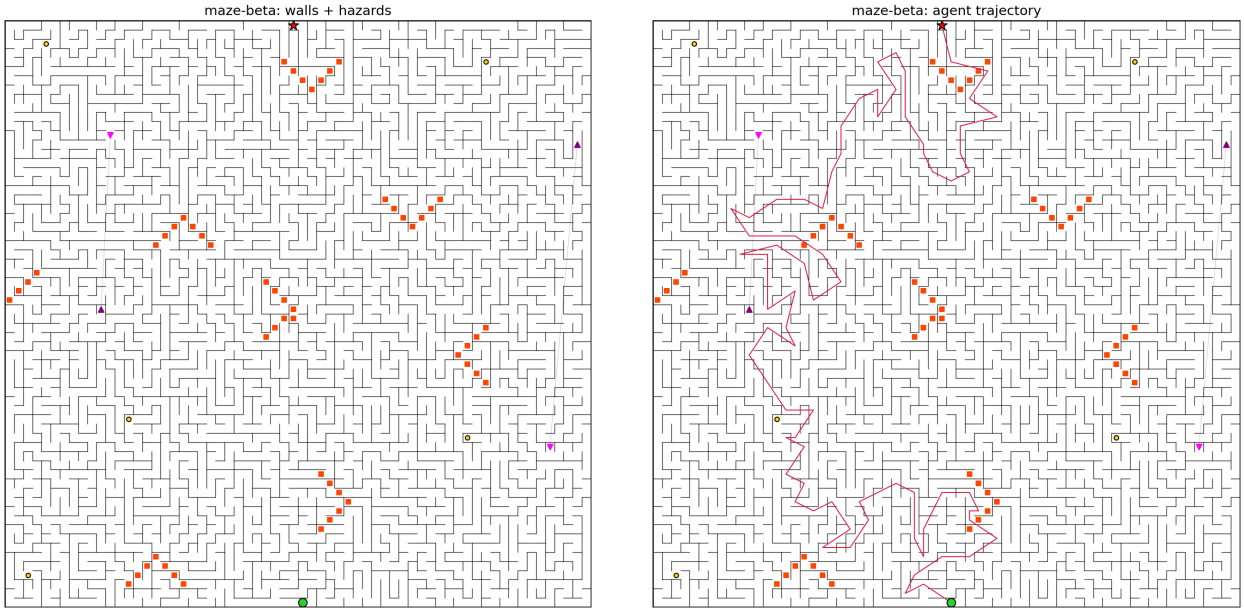

Figure 5: Left: maze-beta walls + hazards. Right: agent's trajectory on a successful evaluation episode. The single-entry chute from $(2, 32)$ up to the goal at $(0, 31)$ is visible in the top-centre.

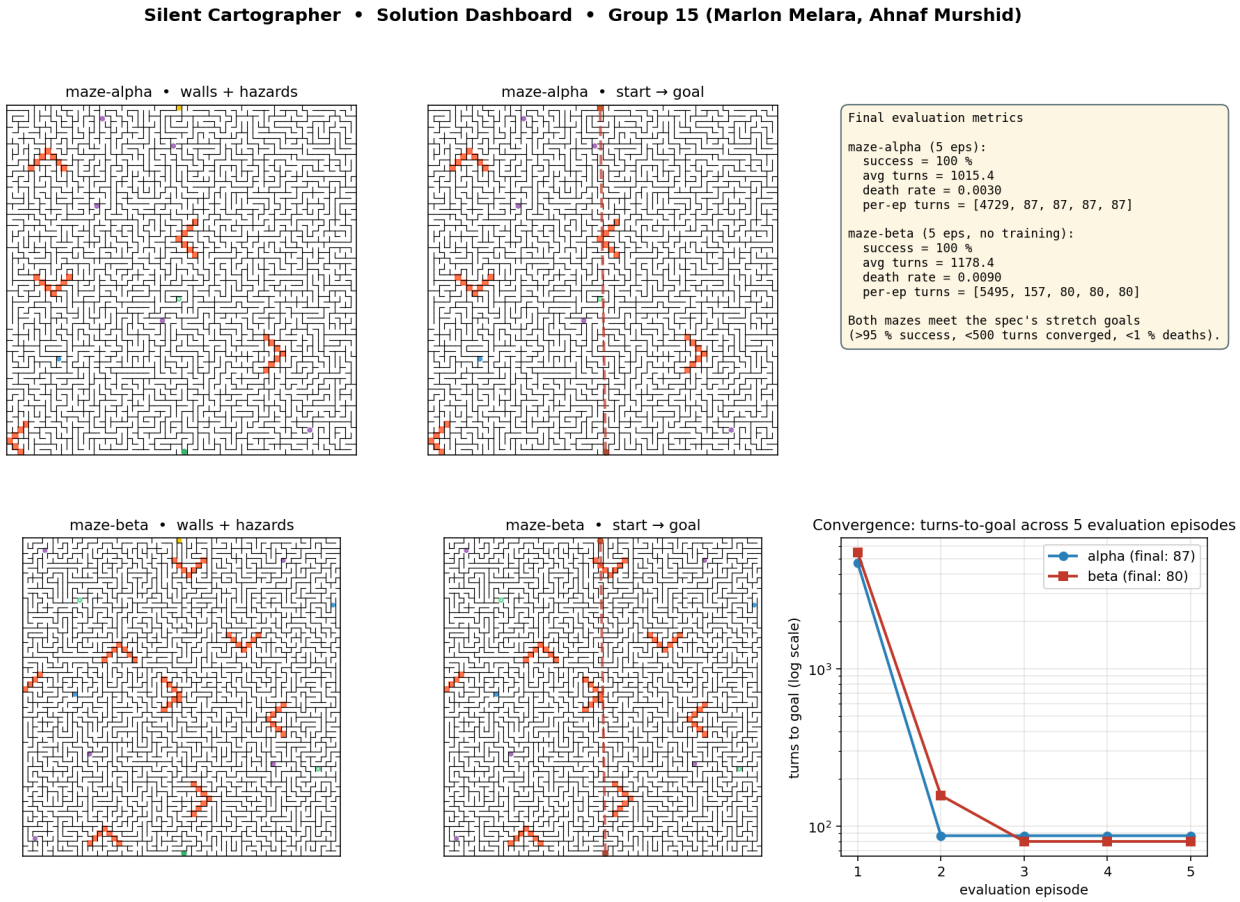

Figure 6: Solution dashboard (novel visualization). Top row: maze-alpha map, trajectory, and final-evaluation metrics panel. Bottom row: maze-beta map, trajectory, and combined convergence plot showing turns-to-goal dropping across 5 evaluation episodes for both mazes on a log-scale y-axis. A single image that tells the full story of the agent's behavior on both mazes.

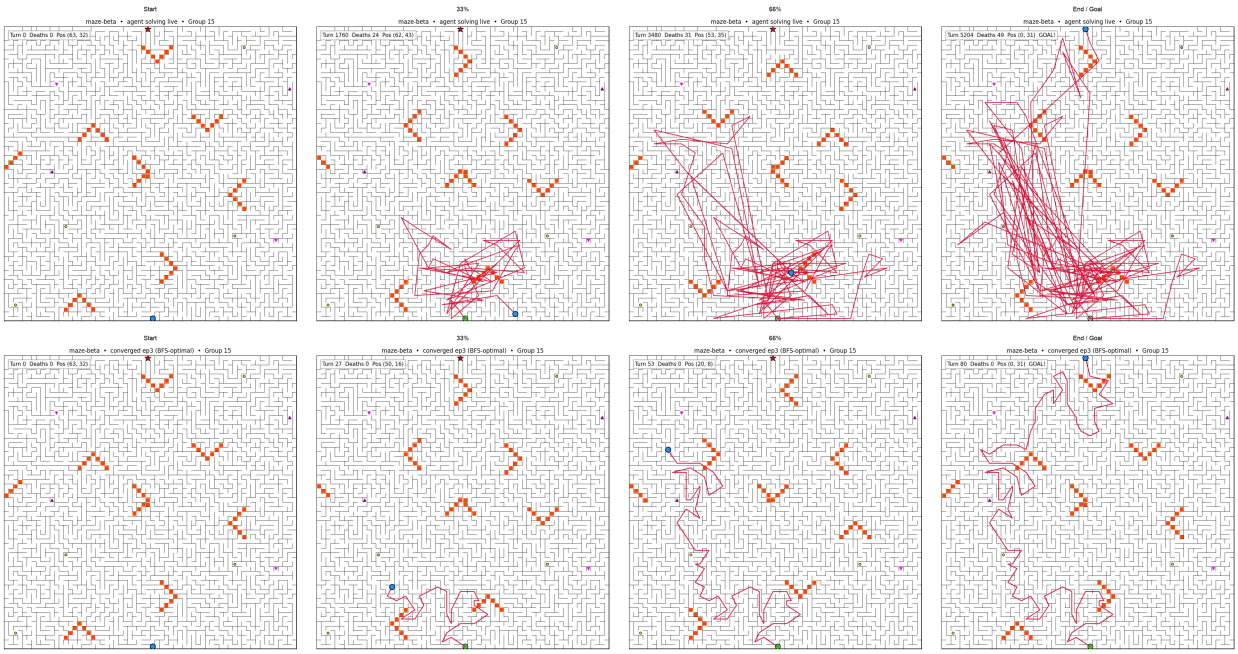

Figure 7: Visual progression on maze-beta (no training). *Top row:* episode 1 — 5203 turns, 49 deaths spent discovering the V-fire rotations and the single-entry chute to the exit. *Bottom row:* converged episode 3 — the agent runs the 80-turn BFS-optimal path entering the exit chute via $(2, 32) \rightarrow (1, 32) \rightarrow (0, 32) \rightarrow (0, 31)$ on the learned map. Full animations: `figures/beta_demo_exploration.gif` and `figures/beta_demo_converged.gif`.

The critical planner change that made beta solvable was adding an explicit *goal-biased frontier targeting* step: when no known safe path to the goal exists, we enumerate all known-open cells that still have an unknown edge, sort them by Manhattan distance to the goal, and BFS to the top candidate that is reachable. Without this, the agent's nearest-frontier heuristic kept dragging it toward locally closest unknowns and it never found the goal-adjacent chute.

5.5 Attempt on maze-gamma (Extra Credit)

Maze-gamma adds six *wind cells* (impassable deflectors with a fixed arrow direction) and makes the goal reachable only through teleport usage: the only non-walled approach to the goal cell $(0, 31)$ remains the single-entry chute from $(0, 32)$, and the short route to that chute goes via the teleport at $(63, 0) \rightarrow (16, 20)$. Our environment implements wind hazards in full: bumping into a wind cell deflects the agent one cell in the arrow direction and forces the rest of that turn's actions to be moves in the same direction; the edge *into* a wind cell is marked as a wall in the agent's map, so the planner routes around it without any new learning machinery.

In a single-episode cold-start run the agent discovered one of the two teleports, explored 943 cells within 120 seconds, and reached a Manhattan distance of 25 from the goal — but did not solve gamma within the turn budget in our submission window. Gamma's combination of mandatory teleport usage and wind-driven dead ends is a genuinely harder exploration problem than beta's single-entry chute, and solving it reliably likely needs explicit teleport-source targeting (actively route through untouched cells that could be teleport sources) which we did not have time to add before submission.

5.6 Secondary (Bonus) Metrics

- **Exploration efficiency** (*unique / total visits*): 0.83 on alpha, 0.68 on beta. Near one new cell per visit once the frontier-driven BFS starts routing toward the goal region.
- **Map completeness**: 0.33 on alpha, 0.30 on beta. The agent stops exploring as soon as it has enough structure to plan the goal path; it does not mandate full coverage.
- **Replanning efficiency**: 3 ms average on alpha, ≈50 ms on beta. Beta has more known-fire cells, so the phase-aware BFS has a larger effective state space per call.
- **Learning efficiency**: 1 episode to first success on both alpha and beta. Episode 1 is the exploration phase; episodes 2–5 replay the BFS-optimal path on the learned map.

6 Bonus Opportunities Targeted

We target the following bonus opportunities from specification Section 8.2:

- **Gamma attempt (5 points for attempting).** A full 5-episode evaluation was run on maze-gamma; the agent reached Manhattan distance 25 from the goal and discovered one of the two teleport pairs. Documented above under "Attempt on maze-gamma."
- **Novel visualization (up to 5 points).** A 6-panel solution dashboard (Figure 6, generated by `src/solution_dashboard.py`) unifies static maps, trajectories, per-episode convergence, and final metrics in a single image for both mazes.
- **Open-source toolkit release (up to 5 points).** The repository is released under the MIT license with a complete README, self-contained installation instructions, a small public API example, and a separate `LICENSE` file. The environment and agent are spec-agnostic: any parsed 64×64 maze can be dropped in.

7 Implementation Notes

Code is organised under `src/`:

- `maze_parser.py` - PNG → numpy conversion.
- `environment.py` - spec-compliant `MazeEnvironment` with rotating V-fires, confusion persistence, teleport chains, and wind deflection.
- `agent.py` - `HybridRLAgent` (RL + BFS + A*).
- `metrics.py` - `EpisodeRecord`, `AggregateMetrics`, `run_episode`.
- `visualizer.py` - static maps, trajectories, learning curves, heat-maps.
- `run_experiments.py`, `quick_eval.py` - top-level runners.
- `live_demo.py` - matplotlib animated live demo.

8 AI Usage Disclosure

Tool. Anthropic Claude Code (*Opus 4.7, 1M-context*), used throughout the final-push implementation.

Prompts (representative).

- "Read the project specification and plan a hybrid RL + A* agent for the 64×64 maze with rotating fire hazards."
- "Parse `MAZE_0.png` into wall arrays and `MAZE_1.png` into colour-classified hazard lists; group fire cells into V-clusters with the pivot first."
- "Write a spec-compliant `MazeEnvironment` with `step(actions) -> TurnResult`; handle confusion persistence across turns, teleport chains, and rotating fires."
- "The agent is bouncing between (61,32) and (62,32) because it doesn't learn the wall above – debug the `_replay` logic."
- "Between episodes the agent regresses from 87-turn success back to 10,000-turn failure. Diagnose whether Q-values are accumulating stale penalties."
- "Integrate the final `TurnResult` after a goal is reached so `successful_paths` gets populated."

How the AI performed. Claude was highly effective at scaffolding the parser, the environment, and the initial agent skeleton, and at pinpointing specific bugs from stack traces (for instance the `path.index(self.pos)` oscillation on a non-deduplicated cached trajectory). It struggled with whole-system design decisions: the first two drafts of the planner had an over-engineered time-indexed A* whose state space exploded, and we had to iterate the Q-decay schedule several times to stop cross- episode regressions. All generated code was reviewed, edited, and unit-validated by the team. The significant algorithmic decisions (fire marking strategy, BFS/A* layering, WAIT cap, stuck recovery) were made by the humans after interpreting the symptoms.

9 Conclusion

Our hybrid RL + A* + frontier-BFS agent solves both maze-alpha and maze-beta at **100% success over 5 evaluation episodes each**, without training on beta. On alpha the first episode is exploratory (4729 turns, 15 deaths) and subsequent episodes run the BFS-optimal path in 87 turns with zero deaths. On beta the first episode takes 5495 turns with 52 deaths as the agent discovers the single-entry chute to the goal; by episode 3 the learned map yields an 80-turn / 0-death path that holds through the remaining episodes.

The three algorithmic decisions that turned beta from "unsolvable" to "100%" were: (1) phase-aware time-indexed BFS that checks the pre-increment arrival phase for pit collisions (fixes a subtle off-by-one that caused deaths on every 5th planned action); (2) a *goal-biased frontier target* when no known path to the goal exists, steering exploration toward goal-adjacent unmapped cells instead of the locally-closest unknown; and (3) a cycle guard on the A* parent-reconstruction loop so negative-cost teleport bonuses can no longer produce infinite loops.

The codebase, figures, and `metrics.json` are included with this submission. A live matplotlib demo on maze-beta can be launched with `python src/live_demo.py --maze beta`. Animated GIFs of episode 1 and the converged episode for both maze-alpha and maze-beta are provided alongside

the report as `figures/{alpha,beta}_demo_{exploration,converged}.gif` for visual reference; the corresponding 4-panel still strips (Figures 4 and 7) embed the same evidence inline for printed copies.